




(a) CE Hessian max eigenvalue decay

(b) t^* vs n_{train} (theory: $t^* \propto n$)

(c) Theorem 5 prediction quality vs Hessian checkpoint



(d) Softmax concentration drives Hessian compression

